



Introduction to PostGIS and using PostGIS from R



What is PostGIS

- Adds support for geographic objects to the PostgreSQL object-relational database
- PostgreSQL already has “geometric types” but native geometries are too limited for GIS data and analysis



What is PostGIS

- PostGIS adds an indexing mechanism to allow queries with spatial restrictions or filters (e.g., “tuples within this bounding box”) to return records very quickly from large tables.



What is PostGIS?

- PostGIS adds a “geometry” data type to the usual data types of the relational database (ie. varchar, integer, date, etc.)
- PostGIS adds spatial predicates and functions using the geometry data type.
 - ST_Distance(geometry, geometry)
 - ST_Area(geometry)
 - ST_Intersects(geometry, geometry))

Geometric Datatypes

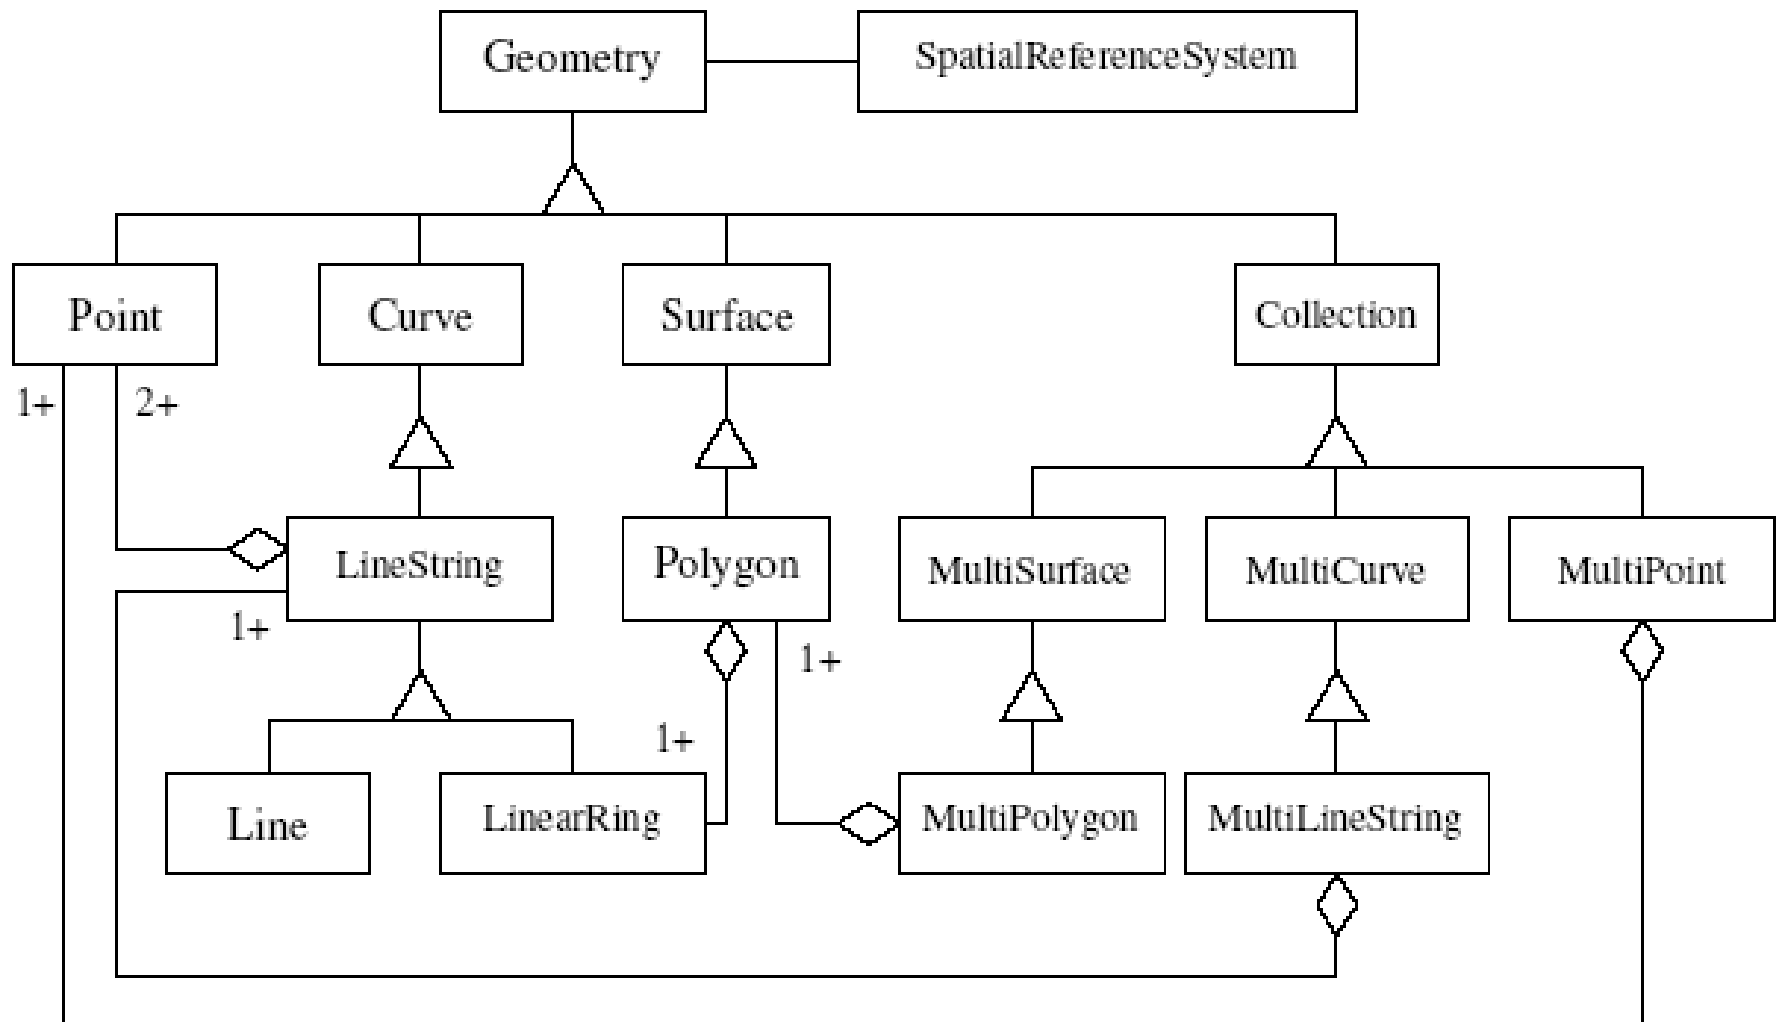


Figure 1: OpenGIS Geometry Class Hierarchy




Geometric Datatypes

- Each geometry has a Well-Known Text representation (WKT)
 - A geometry type
 - A comma-separated list of coordinate pairs
- Internal binary representation
- Serialization to Geography Markup Language (GML) geometry objects

Geometric Datatypes

- Some examples are:


POINT (7 6)



Geometric Datatypes

- Some examples are:


MULTIPOINT (1 4, 1 5, -0.5 5.5314)



Geometric Datatypes

- Some examples are:


`LINESTRING (1 3, 1 1, 3 0)`



Geometric Datatypes

- Some examples are:


MULTILINESTRING ((2 1, 2 2, 3 4),
 (4 3, 3 2, 4 1, 3 3))



Geometric Datatypes

- Some examples are:


POLYGON ((2 5, 3 8, 6 8, 5 4, 2 5),
(3 6, 4 5, 5 7, 4 6, 3 6))



Geometric Datatypes

- Some examples are:

```
MULTIPOLYGON ((( 6 2, 6 4, 8 4, 9 1, 6 2 ),  
                ( 7 3, 8 2, 8 3, 7 3 )),  
                (( 9 2, 8 5, 10 4, 10 3, 9 2 )))
```






2.3 - Geometry Validity

- PostGIS is compliant with the Open Geospatial Consortium's (OGC) OpenGIS Specifications
- Functions require/assume geometries are valid and in many cases, simple.
 - Geometry validity really only applies for areal geometries
 - Geometry simplicity applies to point and linear geometries

Geometry Validity and Simplicity

- By definition, a **LINESTRING** is always *valid*
- It is *simple* if it does not pass through the same point twice.








Geometry Validity

- By definition, a POLYGON is always *simple*.
- It is *valid* if
 - The boundary is made up of *simple* LINESTRINGS
 - boundary rings do not cross
 - holes are completely within the outer ring and touch the outer ring at most one point.
 - it does not contain cutlines / spikes




Geometry Validity

- These POLYGONS are ...



Geometry Validity

- By definition, a MULTIPOLYGON is *valid* iff
 - All elements are valid
 - Element interiors cannot intersect
 - Element boundaries can touch, but only at a finite number of POINTs.





PostgreSQL + PostGIS Installation

- **Windows Installer**

- PostgreSQL + PostGIS
 - Installs itself as a Windows Service
- PgAdmin III
 - Psql Interactive SQL Shell
- <http://postgis.refrations.net/download/windows/>

- **Linux Instaler**

- `sudo apt-get install postgresql-8.2-postgi pgadmin3`
- <http://postgis.refrations.net/docs/ch02.html>

Create Spatial Database

- From the command line

```
createdb -T template_postgis dbname
```

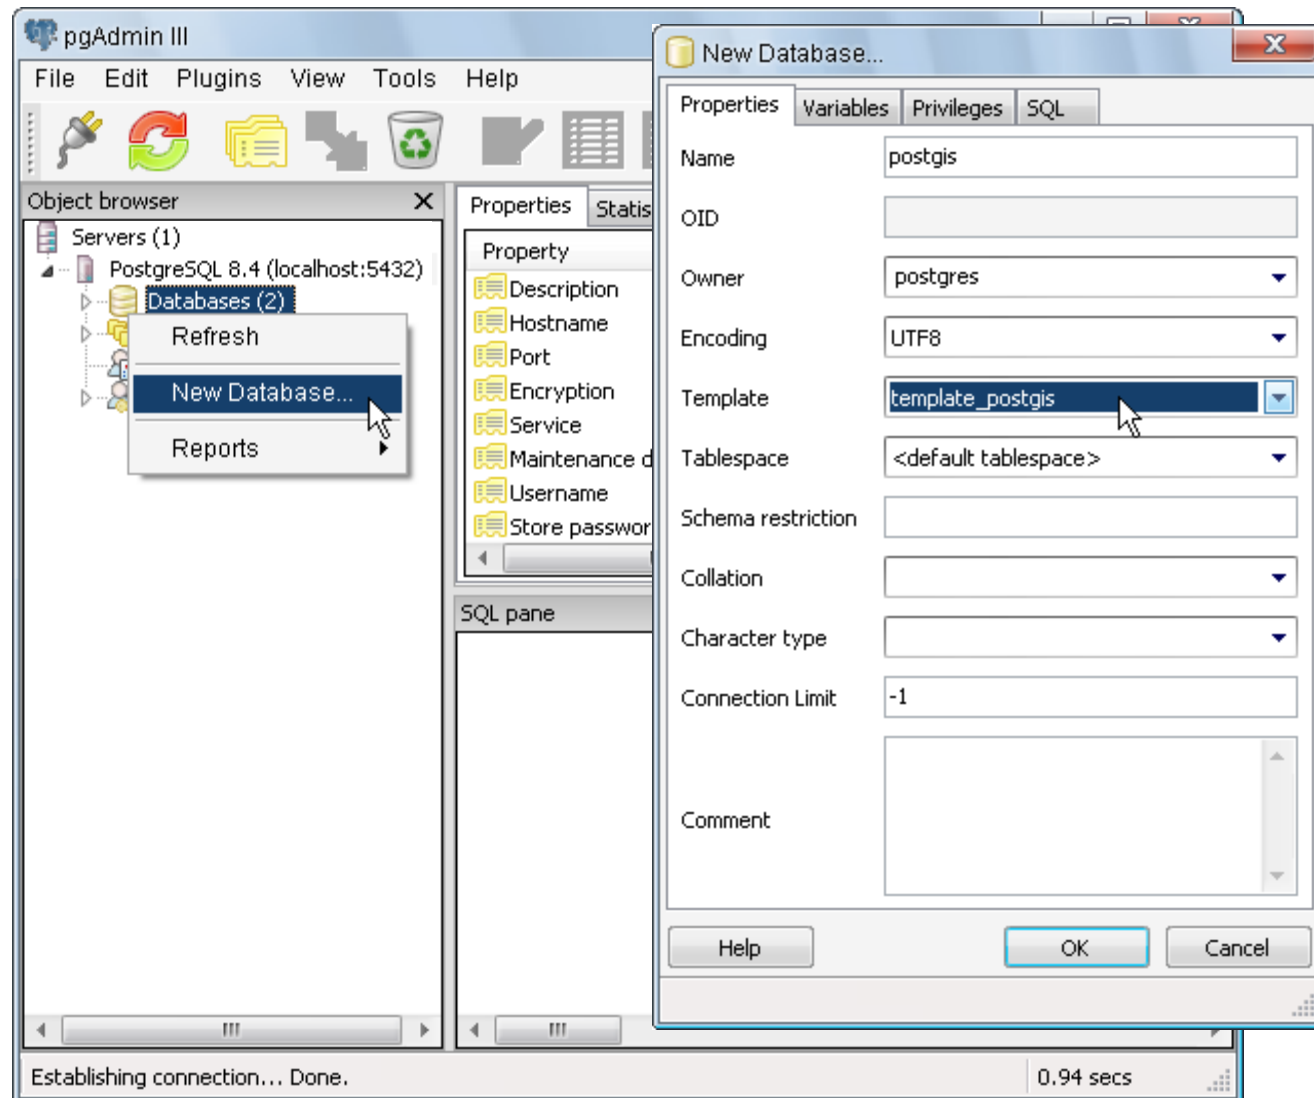
- Create a new database

- Select “template_postgis” as the template

- Verify the installation

- Check for the spatial system tables
 - *spatial_ref_sys* - stores unique coordinate references
 - *geometry_columns* - stores metadata for a spatial column

Create Spatial Database



3

pgAdmin III

File Edit Plugins View Tools Help

Object browser

- Servers (1)
 - PostgreSQL 8.4 (localhost:5432)
 - Databases (3)
 - postgres
 - Schemas (1)
 - public
 - Domains (0)
 - FTS Configurations (0)
 - FTS Dictionaries (0)
 - FTS Parsers (0)
 - FTS Templates (0)
 - Functions (780)
 - Sequences (0)
 - Tables (2)
 - geometry_columns
 - spatial_ref_sys
 - Trigger Functions (2)
 - Views (1)
 - Replication (0)
 - postgres
 - template_postgis
 - Tablespaces (2)
 - Group Roles (0)
 - Login Roles (1)

Properties

Property	Value
Name	geometry_columns
OID	17696
Owner	postgres
Tablespace	pg_default
ACL	
Primary key	f_table_catalog, f_table_schema, f_table_name, f_geometry_column
Rows (estimated)	0
Fill factor	
Rows (counted)	0
Inherits tables	No
Inherited tables count	0
Has OIDs?	Yes
System table?	No

SQL pane

```
-- Table: geometry_columns
-- DROP TABLE geometry_columns;

CREATE TABLE geometry_columns
(
    f_table_catalog character varying(256) NOT NULL,
    f_table_schema character varying(256) NOT NULL,
    f_table_name character varying(256) NOT NULL,
    f_geometry_column character varying(256) NOT NULL,
    coord_dimension integer NOT NULL,
    srid integer NOT NULL,
    "type" character varying(30) NOT NULL,
    CONSTRAINT geometry_columns_pk PRIMARY KEY (f_table_catalog, f_table_schema, f_table_name, f_geometry_column)
)
WITH (
    OIDS=TRUE
);
```

Retrieving Table details... Done.

0.09 secs

Spatially-Enable an Existing DB

- From the command line

```
psql -f "%PGHOME%\share\contrib\postgis-1.5\postgis.sql" dbname  
psql -f "%PGHOME%\share\contrib\postgis-1.5\spatial_ref_sys.sql" dbname
```

- Using a User Interface

- Connect to your existing database
- Open the “Execute Arbitrary SQL Queries” window
- Load/run the PostGIS extension (postgis.sql)
- Load/run the PostGIS spatial reference systems (spatial_ref_sys.sql)

Simple Spatial SQL

- “Manually” create geometries

```
CREATE TABLE points
```

```
  (pt geometry, name varchar);
```

```
INSERT INTO points VALUES
```

```
  ('POINT(0 0)', 'origin');
```

```
INSERT INTO points VALUES
```

```
  ('POINT(5 0)', 'X Axis');
```

```
INSERT INTO points VALUES
```

```
  ('POINT(0 5)', 'Y Axis');
```

```
SELECT name, ST_AsText(pt),
```

```
       ST_Distance(pt, 'POINT(5 5)')
```

```
FROM points;
```



PostGIS Reference

- **Function References:**

<http://postgis.refrations.net/docs/ch07.html>



Loading Shape Files

- **Shape File (Misnomer! 3+ Files!)**
 - .shp = geometry
 - .dbf = attributes (string, number, date)
 - .shx = utility index
- **PostGIS/PostgreSQL Table**
 - Columns can be geometry
 - Columns can be attributes
- **One Shape File = One PostGIS Table**

Loading Shape Files

- **shp2pgsql [opts] shapefile tablename**
 - **shp2pgsql -i -s 3005 -D bc_pubs.shp
bc_pubs > bc_data.sql**
- **Read in .shp file**
- **Write out .sql file**
- **Load .sql file into PostgreSQL**
 - **using psql**
 - **using PgAdmin**


Creating Spatial Indexes

- An example index creation command is


```
CREATE INDEX points_gidx ON points  
USING GIST (pt);
```

Using Spatial Indexes

- Spatial index operator is “&&”
 - “Bounding Boxes Interact”




A && B = TRUE



A && B = FALSE

Using Spatial Indexes

- Bounding Boxes are not enough!



A && **B** = TRUE

_ST_Intersects(A, B) = FALSE

- Two pass system required
 - Use bounding boxes to reduce candidates
 - Use real topological test to get final answer

Using Spatial Indexes

- Index operations (& &) are built into common functions for automatic use, but you can use them separately if you like.
 - ST_Intersects(G1,G2)
 - $G1 \ \&\& \ G2 \text{ AND } _ST_Intersects(G1,G2)$
 - ST_Contains(G1,G2)
 - ST_Within(G1,G2)
 - ST_Touches(G1,G2)
 - ST_DWithin(G1,G2,D)
 - $G1 \ \&\& \ ST_Expand(G2,D) \text{ AND } ST_Distance(G1,G2) > D$



Spatial Analysis in SQL

- A surprising number of traditional *GIS analysis* questions can be answered using a spatial database and SQL.
- *GIS analysis* is generally about filtering spatial objects with conditions, and summarizing the results - and that is exactly what databases are very good at.



Distance Queries

- It's easy to write inefficient queries for distances in PostGIS because it's not always obvious how to use the index in a distance query.

Distance Queries

Question: How many wildlife habitats are within 20 kilometers of the municipality of Oliver?

First, find out where Oliver is located.

```
SELECT ST_AsText(ST_Centroid(the_geom))  
FROM bc_municipalities  
WHERE name = 'Oliver';
```

```
POINT(1470065.29710885 484600.794443135)
```

Distance Queries

Then, use that location to sum all habitats within 20km. The “obvious” way does not use any indexes.

```
SELECT count(*) AS num_habitats
FROM bc_habitat_areas
WHERE
    ST_Distance(
        the_geom,
        ST_GeomFromText('POINT(1470065 484600)', 3005)
    ) < 20000;
```

Distance Queries

The “optimized” way is to use `ST_DWithin()` which silently adds an index into the mix.

```
SELECT count(*) AS num_habitats
FROM bc_habitat_areas
WHERE
    ST_DWithin(
        the_geom,
        ST_GeomFromText(
            'POINT(1470065 484600)', 3005),
        20000
    );
```

Distance Queries

Here's the definition of `ST_DWithin()`.

```
CREATE FUNCTION ST_DWithin(geometry, geometry,  
float8)  
  RETURNS boolean AS '  
    SELECT  
      $1 && ST_Expand($2,$3) AND  
      $2 && ST_Expand($1,$3) AND  
      ST_Distance($1, $2) < $3  
  ' LANGUAGE 'SQL' IMMUTABLE;
```

Spatial Joins

- **Standard joins use a common key**

```
SELECT a.var1, b.var2  
FROM a, b  
WHERE a.id = b.id
```

- **Spatial joins are based on a spatial relationship**

```
SELECT a.var1, b.var2  
FROM a, b  
WHERE ST_Intersects(a.geom, b.geom)
```


Spatial Joins

- The previous distance query can be expressed using a single spatial join.

```
SELECT count(*) AS num_habitats
FROM
  bc_municipalities a,
  bc_habitat_areas b
WHERE
  ST_DWithin(
    a.the_geom,
    b.the_geom,
    20000)
AND
  a.name ILIKE 'Oliver%';
```

Overlays

- Table-on-table overlays are possible with the `ST_Intersection()` function
 - `ST_Intersects(a,b)` returns `BOOLEAN`
 - `ST_Intersection(a,b)` returns `GEOMETRY`



`ST_Intersects()` = `TRUE`

`ST_Intersection()` =





Overlays

Question: Create a new table of habitat areas clipped by the Chilliwack municipal boundary.

Overlays

```
CREATE TABLE ch_habitat_areas AS
SELECT
    ST_Intersection(h.the_geom, m.the_geom)
        AS intersection_geom,
    ST_Area(h.the_geom) AS va_area,
    h.*,
    m.name AS municipality_name
FROM
    bc_habitat_areas h,
    bc_municipalities m
WHERE ST_Intersects(h.the_geom, m.the_geom)
AND m.name = 'Chilliwack';
```



Coordinate Projection

- Every geometry in PostGIS has a “spatial referencing identifier” or “SRID” attached to it.
- The SRID indicates the spatial reference system the geometry coordinates are in.

Coordinate Projection

- View the SRID of geometries using the ST_SRID() function

```
SELECT ST_SRID(the_geom)
FROM bc_roads LIMIT 1;
```

3005

- What's 3005?

```
SELECT srtext
FROM spatial_ref_sys
WHERE srid = 3005;
```

PROJCS["NAD83 / BC Albers",...

- Ah, it's "BC Albers"

Coordinate Projection

- What's 3005 again?

```
SELECT proj4text  
FROM spatial_ref_sys  
WHERE srid = 3005;
```

```
"+proj=aea +lat_1=50 +lat_2=58.5  
+lat_0=45 +lon_0=-126 +x_0=1000000  
+y_0=0 +ellps=GRS80 +datum=NAD83  
+units=m +no_defs"
```

- PROJ4 is coordinate re-projection library used by PostGIS

Coordinate Projection

- Coordinate Re-projection is done using the ST_Transform() function

```
SELECT ST_AsEWKT(the_geom)
FROM bc_roads
LIMIT 1;
```

```
"SRID=3005;MULTILINESTRING( (
    1004687.04355194 594291.053764096,
    1004729.74799931 594258.821943696,
    1004808.0184134 594223.285878035,
    1004864.93630072 594204.422638658,
    1004900.50302171 594200.005856311
) ) "
```

Extracting Spatial Data

- *One of the utility applications that accompany a typical PostGIS installation is the `pgsql2shp` command-line utility program that can extract a table, view, or any custom query that contains a geometry column into a shapefile.*

USAGE:

```
pgsql2shp [<options>] <database> [<schema>.]<table>
```

```
pgsql2shp [<options>] <database> <query>
```

Basic Queries

How many points were used to digitize the road named “Cartmell Rd”?

```
SELECT ST_NPoints(the_geom)
FROM bc_roads
WHERE name = 'Cartmell Rd';
```

Basic Queries

How many holes or interior rings does the habitat area with hab_id = 891 (“Marbled Murrelet”) have?

```
SELECT ST_NumInteriorRings(the_geom)
FROM bc_habitat_areas
WHERE hab_id = 891;
```


Basic Queries

What is the length in kilometers of all roads named 'Main St'?

```
SELECT sum(ST_Length(the_geom)) / 1000  
       AS total_length  
FROM bc_roads  
WHERE name = 'Main St';
```

44.8489926296202

Basic Queries

What is the total area of all municipalities in hectares?

```
SELECT sum(ST_Area(the_geom)) / 10000  
       AS total_length  
FROM bc_municipalities;
```

1186482.22780945

Basic Queries

What is the average road length (treat every entry in the bc_roads table as a separate road)?

```
SELECT avg(ST_Length(the_geom))  
FROM bc_roads;
```

```
434.230644730819
```

Basic Queries

How many roads are completely within the municipality of 'Hope'?

```
SELECT count(*)  
FROM bc_roads r,  
      bc_municipalities m  
WHERE ST_Within(r.the_geom, m.the_geom)  
AND m.name = 'Hope';
```

Basic Queries

What is the latitude of the most southerly hospital in the province?

```
SELECT min(ST_Y(  
            ST_Transform(the_geom, 4269)  
        )) AS lat  
FROM bc_hospitals;
```

48.4657953714625

Basic Queries

According to the datasets, is there a hospital in “Sicamous”?

```
SELECT count(*)  
FROM bc_municipalities a,  
      bc_hospitals b  
WHERE a.name = 'Sicamous'  
AND ST_Contains(a.the_geom, b.the_geom);
```

0

Basic Queries

What is the name of the closest hospital from Sicamous and how far away is it?

```
SELECT a.name, b.name,  
       ST_Distance(a.the_geom, b.the_geom)  
FROM bc_municipalities a,  
      bc_hospitals b  
WHERE a.name = 'Sicamous'  
ORDER BY ST_Distance ASC  
LIMIT 1;
```

name	name	st_distance
Sicamous	Shuswap Lake General Hospital	22943.88318

Basic Queries


What are the names of the pubs located within 100 meters from “Granville St”, located in downtown Vancouver?

```
SELECT DISTINCT c.name
FROM bc_municipalities a,
      bc_roads b,
      bc_pubs c
WHERE a.name = 'Vancouver'
AND b.name = 'Granville St'
AND ST_Contains(a.the_geom, b.the_geom)
AND ST_DWithin(b.the_geom, c.the_geom, 100)
ORDER BY c.name;
```




RODBC + RGDAL Libraries

- An ODBC library for the R environment.
- <http://cran.r-project.org/web/packages/RODBC/>
- RGDAL, an R Geospatial Data Abstraction Layer
- <http://cran.r-project.org/web/packages/rgdal/>



```
library(RODBC)
db <- odbcConnect('gisdatabase', uid='username', pwd='password')
sql <- 'select GeometryType(the_geom), NumGeometries(the_geom),
      asewkt(the_geom) as asewkt, gid from smr150_region'
res <- sqlQuery(db, sql)

geomtype <- as.character(res$geometrytype)
geomnum <- as.character(res$numgeometries)
geom <- as.character(res$asewkt)
geomdesc <- as.character(res$gid)

res <- sqlQuery(db, "SELECT attr, ST_X(the_geom) AS x,
                  ST_Y(the_geom) AS y FROM yourtable WHERE ...")
coordinates(data_frame) <- ~x+y

plot(coordinates)
```